



**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación

Introduction to OmpSs-2

Rosa M. Badia, Xavier Teruel, Xavier Martorell



Barcelona, September 6th, 2022

Introduction to StarSs

The StarSs “Granularities”

StarSs

OmpSs

COMPSs
PyCOMPSs

@ SMP

@ GPU

@ FPGA

@ Cluster

Average task Granularity:

100 microseconds – 10 milliseconds

1second - 1 day

Address space to compute dependences:

Memory

Files, Objects (SCM)

Language binding:

C, C++, FORTRAN

Java, Python, C/C++

Courses: October '22
March '23

Courses: January '23

Parallel

Ensemble, workflow

The parallel programming revolution

Parallel programming in the past

- Where to place data
- What to run where
- How to communicate

Schedule @ programmers mind

Static

Complexity: Increasing divergence
between our mental model and reality

System variability

Parallel programming in the future

- What do I need to compute
- What data do I need to use
- Hints (not necessarily very precise) on potential concurrency, locality,...

Schedule @ system

Dynamic

BSC Vision in the programming revolution

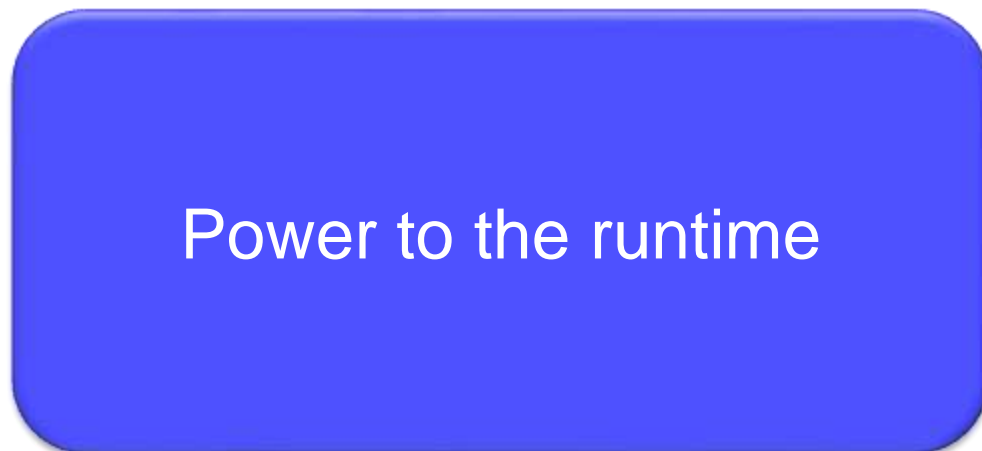
Need to decouple again



Application logic
Arch. independent

PM: High-level, clean, abstract interface





General purpose
Task based
Single address space



“Reuse”
architectural ideas
under
new constraints

ISA / API

Address spaces
(hierarchy, transfers)
Control flow (cores,
threads, accelerators)



OmpSs Concept

The StarSs family of programming models


« Key concept

- **Sequential task based** program on **single address/name space + directionality annotations**
- Happens to execute parallel: Automatic run time computation of dependencies between tasks
 - ~ Sequential equivalence

« StarSs main characteristics


- Dependences: Tasks instantiated but not ready
 - Order IS defined by task creation times
 - Lookahead
 - Avoid stalling the main control flow when a computation depending on previous tasks is reached
 - Possibility to “see” the future searching for further potential concurrency
- Locality aware
- Homogenizing heterogeneity

StarSs: a sequential program ...



```
void Cholesky( float *A[NT][NT] ) {  
  int i, j, k;  
  for (k=0; k<NT; k++) {  
  
    spotrf (A[k*NT+k]) ;  
    for (i=k+1; i<NT; i++) {  
  
      strsm (A[k][k], A[k][i]);  
    }  
    for (i=k+1; i<NT; i++) {  
      for (j=k+1; j<i; j++) {  
  
        sgemm( A[k][i], A[k][j], A[j][i]);  
      }  
  
      ssyrk (A[k][i], A[i][i]);  
    }  
  }  
}
```



StarSs: ... with directionality annotations ...



OmpSs-2 annotations

```
void Cholesky( float *A[NT][NT] ) {  
  int i, j, k;  
  for (k=0; k<NT; k++) {  
    #pragma oss task inout (A[k][k])  
    ● spotrf (A[k][k]) ;  
    for (i=k+1; i<NT; i++) {  
      #pragma oss task in (A[k][k]) inout (A[k][i])  
      ● strsm (A[k][k], A[k][i]);  
    }  
    for (i=k+1; i<NT; i++) {  
      for (j=k+1; j<i; j++) {  
        #pragma oss task in (A[k][i], A[k][j]) inout (A[j][i])  
        ● sgemm( A[k][i], A[k][j], A[j][i]);  
      }  
      #pragma oss task in (A[k][i]) inout (A[i][i])  
      ● ssyrk (A[k][i], A[i][i]);  
    }  
  }  
}
```


StarSs: ... that happens to execute in parallel



OmpSs-2 annotations


```

void Cholesky( float *A[NT][NT] ) {
  int i, j, k;
  for (k=0; k<NT; k++) {
    #pragma oss task inout (A[k][k])
    ● spotrf (A[k][k]) ;
    for (i=k+1; i<NT; i++) {
      #pragma oss task in (A[k][k]) inout (A[k][i])
      ● strsm (A[k][k], A[k][i]);
    }
    for (i=k+1; i<NT; i++) {
      for (j=k+1; j<i; j++) {
        #pragma oss task in (A[k][i], A[k][j]) inout (A[j][i])
        ● sgemm( A[k][i], A[k][j], A[j][i]);
      }
      #pragma oss task in (A[k][i]) inout (A[i][i])
      ● ssyrk (A[k][i], A[i][i]);
    }
  }
}
    
```




Decouple how we write/think (sequential) from how it is executed

StarSs vs OpenMP 3.0 ...



```
void Cholesky( float *A[NT][NT] ) {
  int i, j, k;
  for (k=0; k<NT; k++) {
    ● spotrf (A[k][k]);
    for (i=k+1; i<NT; i++)
      ● #pragma omp task
        ● strsm (A[k][k], A[k][i]);
    #pragma omp taskwait
    for (i=k+1; i<NT; i++) {
      for (j=k+1; j<i; j++)
        ● #pragma omp task
          ● sgemm( A[k][i], A[k][j], A[j][i]);
        ● #pragma omp task
          ● ssyrk (A[k][i], A[i][i]);
        #pragma omp taskwait
    }
  }
}
```

OpenMP annotations:
tasks



OmpSs-like dependences
also in OpenMP since
release 4.0 !!!

```
● spotrf (A[k][k]);
#pragma omp parallel for
for (i=k+1; i<NT; i++)
  ● strsm (A[k][k], A[k][i]);
for (i=k+1; i<NT; i++) {
  #pragma omp parallel for
  for (j=k+1; j<i; j++)
    ● sgemm( A[k][i], A[k][j], A[j][i]);
    ● ssyrk (A[k][i], A[i][i]);
  }
}
```

OpenMP annotations
worksharings

Compiling OmpSs-2 programs

« Compiling

```
mcc --ompss-2 -c bin.c
```

« Linking

```
mcc --ompss-2 -o bin bin.o
```

« Compiler accepts several GCC options

« Compatibility with OpenMP

- --openmp-compatibility
- Accepts OmpSs/OpenMP directives
 - #pragma omp ...
- Do not use in this course. Let the omp directives for the GCC compiler
 - e.g., OMP SIMD

Related Models and environments

« Programming models

- OpenMP 4.5
- Cilk++
- TBB
- OpenACC
- CUDA
- OpenCL
- PGAS models
 - UPC
 - X10
 - Chapel

« Runtimes

- ParalleX, HPX @ LSU
- Swarm
- StarPU @ U. Bordeaux
- PLASMA, ... @ UTK
- Argobots @ ANL
- Open Community Runtime
- ...

Basic OmpSs-2

Execution Model

❧ Thread-pool model


- OpenMP parallel not supported (it was ignored in OmpSs-1)

❧ All threads created on startup

- One of them starts executing main (on SMP device)
- P-1 workers execute SMP tasks
 - `$ taskset -c 0-$(P-1)`
 - Equivalent to the OpenMP `OMP_NUM_THREADS`
- One representative (OpenCL/CUDA/FPGA) per device/accelerator

❧ All get work from a task pool

- And can generate new work
- Work is labeled with possible “targets”
- Single or multiple work-queues



Memory Model

« From the point of view of the programmer a single naming space exists

- The standard sequential/shared memory address space



« From the point of view of the runtime and target platform, different possible scenarios

- Pure SMP:
 - Single address space
- Distributed/heterogeneous (cluster, GPUs, ...):
 - Multiple address spaces exist
 - Versions of same data may exist in multiple of these address spaces
 - Data consistency ensured by the implementation

Main element: tasks

Task

- Computation unit. Amount of work (granularity) may vary in a wide range (μ secs to msecs or even seconds), may depend on input arguments,...
- Once started can execute to completion independent of other tasks
- Can be declared inlined or outlined

States:

- **Instantiated**: when task is created. Dependences are computed at the moment of instantiation. At that point in time a task may or may not be ready for execution
- **Ready**: When all its input dependences are satisfied, typically as a result of the completion of other tasks
- **Active**: the task has been scheduled to a processing element. Will take a finite amount of time to execute.
- **Completed**: the task terminates, its state transformations are guaranteed to be globally visible and frees its output dependences to other tasks.

OmpSs main initial contribution: Dependences

```
#pragma oss task in(lvalue_expr-list) \  
                out(lvalue_expr-list) \  
                inout(lvalue_expr-list)
```

- ⌘ Specify data that has to be available before activating task
- ⌘ Specify data this task produces that might activate other tasks
- ⌘ Contiguous / array regions
- ⌘ You do NOT specify dependences but information for the runtime to compute them
- ⌘ Contributed to OpenMP 4.0

⌘ Relaxations:

- No need to specify for ALL data read/written, can use representatives

- Order in inout chains

- Programmer responsible of possible races

```
#pragma oss concurrent (...)
```

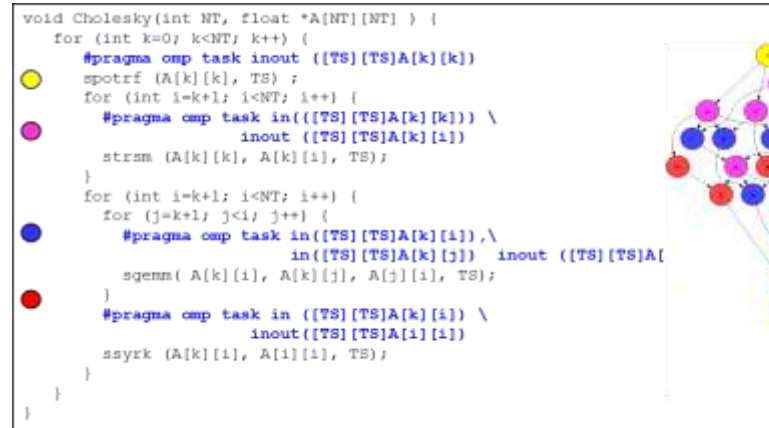
```
#pragma oss critical | atomic
```

```
#pragma oss commutative (...)
```

Inlined and outlined tasks

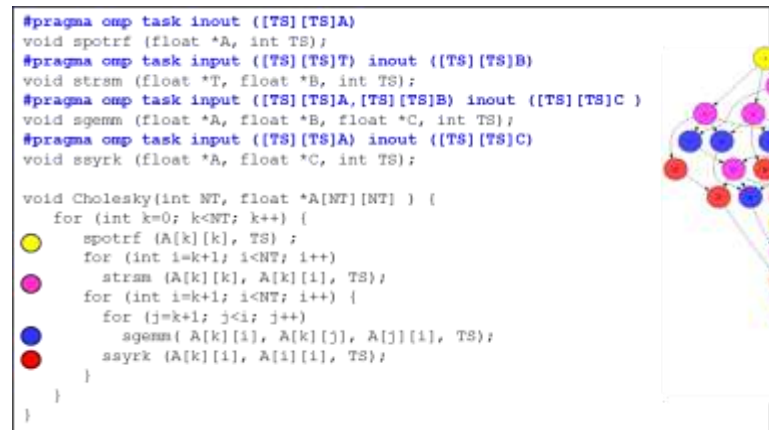
Pragmas inlined

- Pragma applies to immediately following statement
- The compiler outlines the statement (as in OpenMP)



Pragmas outlined

- Attached to function declaration
 - All function invocations become a task
 - The programmer gives a name, this enables to later provide several implementations



Inlined directives

Inlined task instantiation and wait

`#pragma oss task`

- Instantiates a task and control flow continues

`#pragma oss taskwait`

- Suspends the current control flow until all children tasks are completed

```
int  main  (    )
{
    int X[100];

    #pragma  oss  task
    for (int i=0; i< 100; i++) X[i]=i;
    #pragma  oss  taskwait

    ...
}
```



Inlined task instantiation and wait

`#pragma oss task`


- Instantiates a task and control flow continues

`#pragma oss taskwait`

- Suspends the current control flow until all children tasks are completed

```
void traverse_list ( List l )
{
    Element e ;
    for ( e = l-> first; e ; e = e->next ) {
        #pragma oss task
        process ( e ) ;
    }

    #pragma oss taskwait
}
```



Without taskwait the subroutine will return immediately after spawning the tasks allowing the calling function to continue spawning tasks

Data scope

“ The data accessed by a task may be (as in OpenMP):

- **Shared:** the task uses the original variable in the context where it is instantiated
- **Private:** the task uses an uninitialized private copy of the data allocated at task activation time
- **Firstprivate:** the task uses a private copy of the data initialized to the value of the variable in the instantiating context at task instantiation time
- **NO thread private scope.**

“ OmpSs scoping rules == OpenMP scoping rules

Dependences

```
#pragma oss task in(lvalue_expr-list) \  
                out(lvalue_expr-list) \  
                inout(lvalue_expr-list)
```

- ⌘ Specify data that has to be **available** before activating task
- ⌘ Specify data this task produces that might activate other tasks
- ⌘ Contiguous / array sections
- ⌘ You do NOT specify dependencies but information for the runtime to compute them

⌘ Relaxations:

- No need to specify for ALL data read/written, can use representatives
- Order of inout chains
 - Programmer responsible of possible races

```
#pragma oss concurrent (...)
```

```
#pragma oss critical | atomic
```

```
#pragma oss commutative (...)
```


Defining dependences for inlined tasks


❧ Clauses that express data direction:

- In, out, inout
- These clauses **imply shared scope** for the variable referenced in them
- The argument is an **lvalue** expression referring to the object that will be used to compute dependences
- Uses of the lvalue expression in the task body reference the global object, whose value is guaranteed to have been produced by the “latest” task to “modify” it (latest out/inout task in sequential program order) if any.

❧ Dependences computed at runtime taking into account these clauses

- Flow, anti, output

```
#pragma oss task out( x )  
x = 5;                               //1  
#pragma oss task in( x )  
printf("%d\n" , x ) ;                 //2  
#pragma oss task inout( x )  
x++;                                  //3  
#pragma oss task in( x )  
printf ("%d\n" , x ) ;                 //4
```



OpenMP: Dependencies

- « Dependencies are now supported in OpenMP 4.0/4.5/5.0/5.1
 - Different syntax but same semantics

```
#pragma oss task [ depend (in: var_list) ] [ depend(out: var_list) ] [ depend(inout: var_list) ]  
{  
    // code block  
}  
  
//      var_list: var_name | var_name[start:size]...  
  
//      OpenMP inoutset (concurrent) and mutexinoutset (commutative) not supported yet
```


“OpenMP Application Program Interface. Version 4.0 “ July 2013

and

“OpenMP Application Program Interface. Version 4.5” November 2015

Defining dependences for inlined tasks in OpenMP 4.0

```
#pragma omp task depend (out:x )  
x = 5;                               //1  
#pragma omp task depend (in:x)  
printf("%d\n" , x ) ;                 //2  
#pragma omp task depend (inout:x)  
x++;                                  //3  
#pragma omp task depend (in:x )  
printf ("%d\n" , x ) ;                //4
```




Partial control flow synchronization (OmpSs)

`#pragma taskwait on (lvalue_expr-list)`

- Expressions allowed are the same as for the directionality clauses
- Stalls the encountering control flow until the referenced data is available

```
double A[N][N], B[N][N], C[N][N], D[N][N], E[N][N],  
       F[N][N], G[N][N], H[N][N], I[N][N], J[N][N];
```

```
main() {  
    #pragma oss task in (A, B) out(C)  
    dgemm(A, B, C);    //1  
    #pragma oss task in (D, E) out(F)  
    dgemm(D, E, F);    //2  
    #pragma oss task in (C, F) out(G)  
    dgemm(C, F, G);    //3  
    #pragma oss task in (A, D) out(H)  
    dgemm(A, D, H);    //4  
    #pragma oss task in (C, H) out(I)  
    dgemm(C, H, I);    //5  
  
    #pragma oss taskwait on (F)  
    printf("result F = %f\n", F[0][0]);  
  
    #pragma oss task in (C, H) out(J)  
    dgemm(H, C, J);    //6  
  
    #pragma oss taskwait  
    printf("result J = %f\n", J[0][0]);  
}
```




Partial control flow synchronization (OmpSs & OpenMP)

`#pragma taskwait on (lvalue_expr-list)`

- Expressions allowed are the same as for the directionality clauses
- Stalls the encountering control flow until the referenced data is available

```
double A[N][N], B[N][N], C[N][N], D[N][N], E[N][N],  
       F[N][N], G[N][N], H[N][N], I[N][N], J[N][N];
```

```
main() {  
    #pragma oss task in (A, B) out(C)  
    dgemm(A, B, C);    //1  
    #pragma oss task in (D, E) out(F)  
    dgemm(D, E, F);    //2  
    #pragma oss task in (C, F) out(G)  
    dgemm(C, F, G);    //3  
    #pragma oss task in (A, D) out(H)  
    dgemm(A, D, H);    //4  
    #pragma oss task in (C, H) out(I)  
    dgemm(C, H, I);    //5  
  
    #pragma oss task in(F)  
    printf("result F = %f\n", F[0][0]);  
  
    #pragma oss task in (C, H) out(J)  
    dgemm(H, C, J);    //6  
  
    #pragma oss taskwait  
    printf("result J = %f\n", J[0][0]);  
}
```



Array sections

« Indicating as in/out/inout subregions of a larger structure:

in (A[i])

→ the input argument is element i of A

« Indicating an array section:

in ([BS]A)

→ the input argument is a block of size BS from address A

OmpSs

in (A[i;BS])

→ the input argument is a block of size BS from address $\&A[i]$

→ the lower bound can be omitted (default is 0)

OmpSs

in (A[i:j])

→ the input argument is a block from element $A[i]$ to element $A[j]$ (included)

→ $A[i:i+BS-1]$ equivalent to $A[i; BS]$

→ the upper bound can be omitted if size is known to the compiler
(default is $N-1$, being N the size)

OmpSs

Array sections

```
int a[N];  
#pragma oss task in(a)
```

=

```
int a[N];  
#pragma oss task in(a[0:N-1])  
//whole array used to compute dependences
```

=

```
int a[N];  
#pragma oss task in(a[0:N])  
//whole array used to compute dependences
```

```
int a[N];  
#pragma oss task in(a[0:3])  
//first 4 elements of the array used to compute dependences
```

```
int a[N];  
#pragma oss task in(a[2:3])  
//elements 2 and 3 of the array used to compute dependences
```

=

```
int a[N];  
#pragma oss task in(a[2:2])  
//elements 2 and 3 of the array used to compute dependences
```

Array sections

```
int a[N][M];  
#pragma oss task in(a[0:N-1][0:M-1])  
//whole matrix used to compute dependences
```


=

```
int a[N][M];  
#pragma oss task in(a[0:N][0:M])  
//whole matrix used to compute dependences
```


```
int a[N][M];  
#pragma oss task in(a[2:3][3:4])  
// 2 x 2 subblock of a at a[2][3]
```

=


```
int a[N][M];  
#pragma oss task in(a[2;2][3;2])  
// 2 x 2 subblock of a at a[2][3]
```



```
int a[N][M];  
#pragma oss task in(a[1:2][0:M-1])  
//rows 1 and 2
```



Array sections




```
for (int j; j<N; j+=BS){  
    actual_size = (N- j> BS ? BS: N-j);  
    #pragma oss task in (vec[j;actual_size]) inout(c[j;actual_size])  
    for (int count = 0; count < actual_size; count ++, j++)  
        c[j] += vec [j] ;  
}
```

dynamic size of
argument

Array sections

```
void matmul(double A[N][N], double B[N][N],
            double C[N][N], unsigned long BS)
{
    int i, j, k;

    for (i = 0; i < BS; i++)
        for (j = 0; j < BS; j++)
            for (k = 0; k < BS; k++)
                C[i][j] += A[i][k] * B[k][j];
}
```




```
void compute(unsigned long NB, unsigned long BS,
             double A[N][N], double B[N][N], double C[N][N])
{
    unsigned i, j, k;

    for (i = 0; i < N; i += BS)
        for (j = 0; j < N; j += BS)
            for (k = 0; k < N; k += BS) {
                #pragma oss task in(A[i;BS][k;BS], B[k;BS][j;BS]) inout(C[i;BS][j;BS])
                matmul (&A[i][k], &B[k][j], &C[i][j], BS);
            }
}
```

Incomplete directionalities specification

- ❧ Directionality not required for all arguments
- ❧ May even refer to variables not used in the body of the task
 - Used to force dependences under complex structures (graphs, ...)

```
main () {  
    int sentinel;  
  
    #pragma oss task out (sentinel)  
    foo (...);  
  
    #pragma oss task in (sentinel)  
    bar (...)  
}
```



Sentinel

- Mechanism to handle complex dependences
 - When difficult to specify proper input/output clauses
- To be avoided if possible
 - The use of an element or group of elements as sentinels to represent a larger data-structure is valid
 - However might make code non-portable to heterogeneous platforms if copy in/out clauses cannot properly specify the address space that should be accessible in the devices

OmpSs Cholesky with OpenMP 4.0 syntax

```
void cholesky_blocked(const int ts, const int nt, double* Ah[nt][nt])
{
    for (int k = 0; k < nt; k++) {
        // Diagonal Block factorization
        #pragma oss task depend( inout: Ah[k][k] ) firstprivate(k,ts) label (potrf)
        omp_potrf (Ah[k][k], ts, ts);
        // Triangular systems
        for (int i = k + 1; i < nt; i++) {
            #pragma oss task depend(in: Ah[k][k] ) depend(inout: Ah[k][i]) firstprivate(k,i,ts) \
                label (trsm)

            omp_trsm (Ah[k][k], Ah[k][i], ts, ts);
        }
        // Update trailing matrix
        for (int i = k + 1; i < nt; i++) {
            for (int j = k + 1; j < i; j++) {
                #pragma oss task depend(in: Ah[k][i], Ah[k][j] ) depend(inout: Ah[j][i] ) \
                    firstprivate(k,i,j,ts) label (gemm)

                omp_gemm (Ah[k][i], Ah[k][j], Ah[j][i], ts, ts);
            }
            #pragma oss task depend(in: Ah[k][i]) depend(inout: Ah[i][i]) firstprivate(k,i,ts) \
                label (syrk)

            omp_syrk (Ah[k][i], Ah[i][i], ts, ts);
        }
    }
    #pragma oss taskwait
}
```

OmpSs: syntax in Fortran

```
!$OSS TASK [IN (...)] [OUT(...)] [INOUT (...)] &  
    [CONCURRENT (...)] [COMMUTATIVE (...)] &  
    [PRIORITY (...)] [LABEL (...)] &  
    [SHARED (...)] [PRIVATE (...)] [FIRSTPRIVATE (...)] [DEFAULT (...)] &  
    [UNTIED] [IF (...)] [FINAL (...)] &  
    [DEVICE ({SMP|CUDA|OPENCL|MPI})] &  
    [NDRANGE (...)] &  
    [IMPLEMENTS (function_name)]* &  
    [SHMEM (...)]  
{ code or function }
```

```
!$OSS TASKWAIT [ON (...)] [NOFLUSH]
```

- Pragmas inlined

```
program example
parameter(N=2048)
integer, parameter :: BSIZE = 64
real v(N), vx(N), vy(N), vz(N)
integer :: jj
...
interface
    subroutine v_mod(BSIZE, v, vx, vy, vz)
        implicit none
        integer, intent(in) :: BSIZE
        real, intent(out) :: v(BSIZE)
        real, intent(in), dimension(BSIZE) :: vx, vy, vz
    end subroutine
end interface
...
do jj=1, N, BSIZE
    !$oss task out (v(jj)) in (vx(jj), vy(jj), vz(jj)) firstprivate (jj) \
                                                label(vmod)

        call v_mod(BSIZE, v(jj), vx(jj), vy(jj), vz(jj))
    !$oss end task
enddo
...
!$oss taskwait
```


```
#pragma omp task concurrent (lvalue_expr) // inoutset in OMP
```

⌘ Relaxed inout directionality clause

- Enables the scheduler to change the order (or **even execute concurrently**) the tasks within the inout chains built by the concurrent clause.
 - Dependences resulting from the regular in, out and inout clauses towards and from the concurrent chain are honoured for **all** tasks in the concurrent chain.
- The programmer assumes responsibility to handle potential races within the concurrent chain using if needed OpenMP pragmas

```
#pragma omp atomic  
#pragma omp critical
```


Serialized reduction pattern



```
for (int j=0; j<N; j+=BS){  
    actual_size = (N- j> BS ? BS: N-j);  
  
    #pragma oss task in (vec[j;actual_size]) inout(result)  
    for (int count = 0; count < actual_size; count ++, j++)  
        result += vec [j] ;  
}  
#pragma oss task input (result)  
    printf ("TOTAL is %d\n", result);  
#pragma oss taskwait
```

Serialization

Concurrent



```
double vec[N];
double result;

for (int j; j<N; j+=BS){
    actual_size = (N- j> BS ? BS: N-j);

    #pragma oss task in (vec[j;actual_size]) concurrent(result)
    { double local_result=0.0;
      for (int count = 0; count < actual_size; count ++){
          local_result += vec [j++] ;
          #pragma oss atomic
          result += local_result;
      }
    }

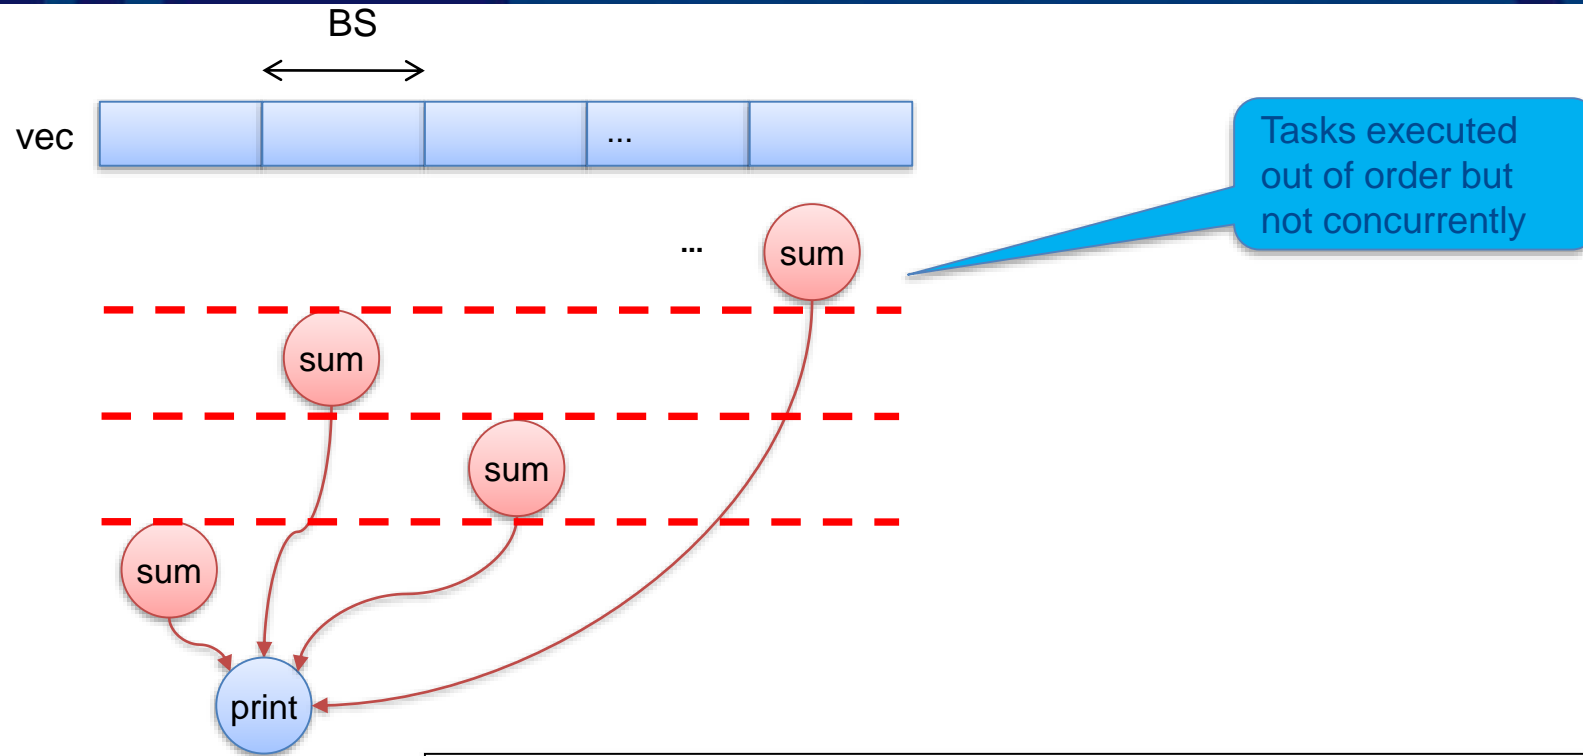
    #pragma oss task input (result)
    printf ("TOTAL is %d\n", result);
    #pragma oss taskwait
```

```
#pragma omp task commutative (lvalue_expr)  
                                // mutexinoutset in OMP
```

⌘ Relaxed inout directionality clause

- Enables the scheduler to change the order, but **not to execute concurrently** the tasks within the inout chains built by the commutative clause.
 - Dependences resulting from the regular in, out and inout clauses towards and from the commutative chain are honoured for all tasks in the commutative chain.

Commutative




No mutual exclusion required




```
for (int j; j<N; j+=BS){  
    actual_size = (N- j> BS ? BS: N-j);  
  
    #pragma oss task in (vec[j;actual_size]) commutative(result)  
    for (int count = 0; count < actual_size; count ++, j++)  
        result += vec [j] ;  
}  
#pragma oss task input (result)  
printf ("TOTAL is %d\n", result);  
#pragma oss taskwait
```

Differences between concurrent and commutative

Tasks timeline: views at same time scale



Histogram of tasks duration: at same control scale



In this case, concurrent is more efficient

... but tasks have more duration and variability

Task nesting

Hierarchical task graph

« Nesting

- As in OpenMP
- Tasks can generate tasks themselves
- Everything said so far now applies within the context of a task whose sequential control flow now becomes the main in the local context

« Hierarchical task dependences

- The generation of children only takes place once father is activated
- Dependences only checked between siblings
- Every father has a separate task graph (Hierarchy)
- There is no implicit taskwait at the end of a task instantiating children

« Impact

- Allows for each level to contribute to express potential parallelism
- Effectively parallelizes task creation and dependence handling


Nesting

```
int Y[4]={1,2,3,4}

int main( )
{
    int X[4]={5,6,7,8};

    for (int i=0; i<2; i++) {
        #pragma oss task out(Y[i]) firstprivate(i,X)
        {
            for (int j=0 ; j<3; j++) {
                #pragma oss task inout(X[j])
                X[j]=f(X[j], j);
                #pragma oss task in (X[j]) inout (Y[i])
                Y[i] +=g(X[j]);
            }
            #pragma oss taskwait
        }
    }
    #pragma oss task inout(Y[0;2])
    for (int i=0; i<2; i++) Y[i] += h(Y[i]);
    #pragma oss task inout (Y[3])
    for (int i=1; i<N; i++) Y[3]=h(Y[3]);

    #pragma oss taskwait
}
```



Nesting

« Top down


- Every level contributes

« Flattening dependence graph

- Increase concurrency
- Take out runtime overhead from critical path

« Granularity control

- final clauses, runtime



Nesting

« Top down


- Every level contributes

« Flattening dependence graph

- Increase concurrency
- Take out runtime overhead from critical path

« Granularity control

- final clauses, runtime



Nesting

Top down


- Every level contributes

Flattening dependence graph

- Increase concurrency
- Take out runtime overhead from critical path

Granularity control

- final clauses, runtime




OmpSs-2: Improved Nesting + Dependences

```
#pragma oss task weakout (a, b, c)
```

```
{  
  #pragma oss task out(a)  
  {...}  
  #pragma oss task out(b)  
  {...}  
  #pragma oss task in(a, b) out(c)  
  {...}  
}
```

After finishing, release
dependences incrementally

Connect Inner
Dependencies
Outwards




```
#pragma oss task weakin(a, c)
```

```
{  
  #pragma oss task in(a)  
  {...}  
  #pragma oss task in(c)  
  {...}  
}
```

Do not wait for these
dependences


Wait only for the inner
ones

Connect Outer
Dependencies
Inwards



OmpSs-2: Improved Nesting + Dependences

#pragma oss task **weakout**(a, b, c)



#pragma oss task **weakin**(a, c)

Our team

- « Jesus Labarta
- « Eduard Ayguade
- « Rosa M. Badia
- « Xavier Martorell
- « Vicenç Bertran
- « Antonio Peña
- « Daniel Jiménez
- « Carlos Álvarez
- « Xavier Teruel
- « Roger Ferrer
- « Marta Garcia
- « Victor Lopez
- « Julian Morillo
- « Kevin Sala
- « Antoni Navarro
- « Antonio Filgueras
- « Daniel Peyrolón
- « Miquel Vidal
- « Orestis Korakitis
- « Simon Garcia
- « Judit Gimenez
- « German Llort
- « Harald Servat (Intel)
- « Aymar Rodríguez
- « Jaume Bosch
- « Guray Ozen (Nvidia/PGI)
- « Artem Cherkashin
- « Marc Josep
- « Marc Marí
- « Ferran Pallarès
- « Albert Navarro (Semidynamics)
- « Alex Duran (Intel)
- « Josep M. Perez
- « Javier Bueno (Metempsy)
- « Judit Planas (Lausanne)
- « Guillermo Miranda (UPC)
- « Florentino Sainz (BBVA)
- « Omer Subasi
- « Javier Arias
- « Sergi Mateo (MongoDB)
- « ...

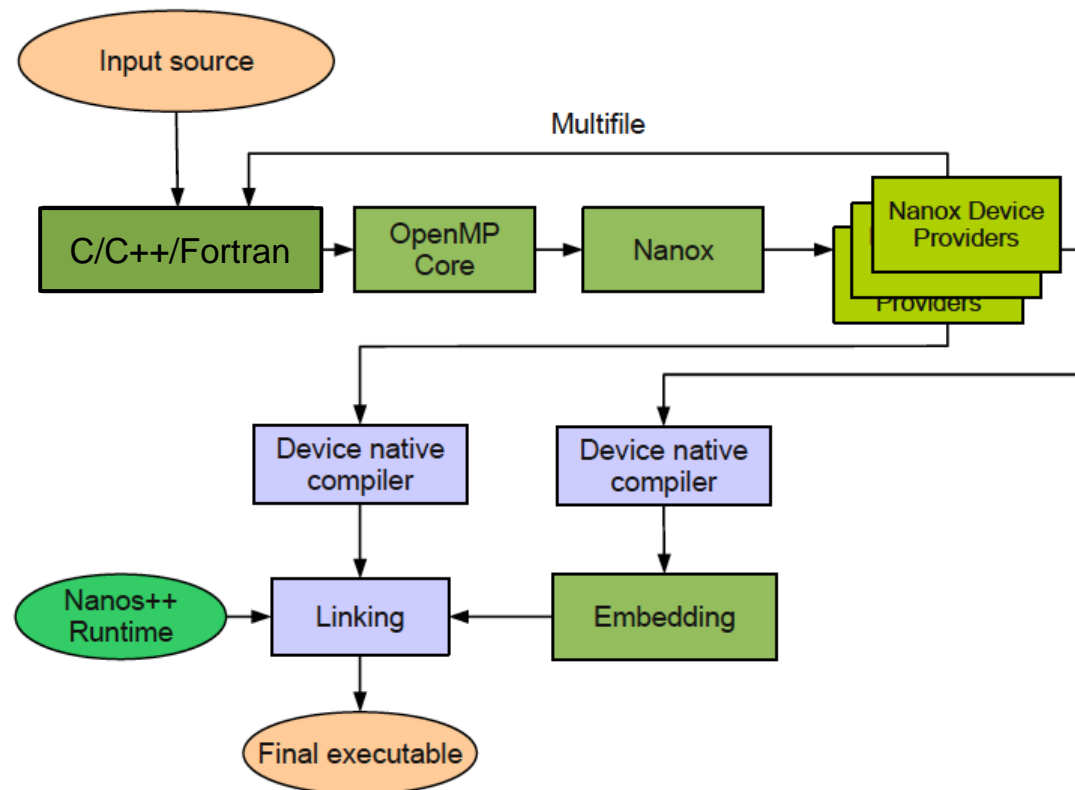
OmpSs environment

Mercurium Compiler

- Recognizes constructs and transforms them to calls to the runtime
- Manages code restructuring for different target devices



- Device-specific handlers
- May generate code in a separate file
- Invokes different back-end compilers
 - gcc, icc, xlc... for regular code
 - nvcc for NVIDIA



❧ Compiling

```
frontend --ompss-2 -c bin.c
```

❧ Linking

```
frontend --ompss-2 -o bin bin.o
```

❧ where *frontend* can be:

mcc	C (gcc backend)
mcxx	C++
mfc	Fortran (gfortran backend)
imcc	C (icc backend)
imcxx	C++ (icpc backend)
imfc	Fortran (ifort backend)

« Compatibility with OpenMP

- `--openmp-compatibility`
- Accepts OmpSs/OpenMP directives
 - `#pragma omp ...`
- Do not use in this course. Let the omp directives for the GCC compiler
- [Useful for the taskloop support
 - From previous editions of the course
 - Now taskloop is already implemented in OmpSs-2]

Compatibility flags:

- -l, -g, -L, -I, -E, -D, -W

Other compilation flags:

-k	Keep intermediate files
--debug	Use Nanos++ debug version
--version	Show Mercurium version number
--verbose	Enable Mercurium verbose output
--Wp,flags	Pass flags to preprocessor (comma separated)
--Wn,flags	Pass flags to native compiler (comma separated)
--Wl,flags	Pass flags to linker (comma separated)
--help	To see many more options :-)

Executing

« Basic execution:

```
> ./bin      # will use all cores/hwthreads available
```

– No LD_LIBRARY_PATH or LD_PRELOAD needed

« Number of threads can be adjusted with “taskset”

```
> taskset -c 20-23 ./bin
```

« In the hands-on

```
> sbatch run-once.sh      # and slurm sets the cpu mask
```

« NANOS6 utility to

- List compiler and compilation options used for the runtime
 - `nanos6-info --runtime-compiler`
 - `nanos6-info --runtime-compiler-flags`
- List runtime details:

```
nanos6-info --runtime-details
```

Runtime path /apps/PM/ompss-2/2020.11.1/lib/libnanos6-optimized.so.0.0.0

Runtime Version 2.5.1 2020-12-22 18:47:30 +0100 443f6d83

Runtime Branch

Runtime Compiler Version g++ (GCC) 6.4.0

Runtime Compiler Flags -DNDEBUG -std=gnu++11 -Wall -Wextra -Wdisabled-optimization -Wshadow -fvisibility=hidden -O3 -fno

Initial CPU List 0-47

NUMA Node 0 CPU List 0-23

NUMA Node 1 CPU List 24-47

Scheduler priority

Dependency Implementation linear-regions-fragmented

Threading Model pthreads

Tracing

❧ Compile and link normally

```
mcc --ompss-2 -c bin.c
```

```
mcc -o bin --ompss-2 bin.o
```

❧ When executing specify which instrumentation module to use:

```
NANOS6_CONFIG_FILE=extrae.xml
```

```
NANOS6_CONFIG_OVERRIDE="version.instrumentation=extrae" ./bin
```

❧ Will generate trace files in executing directory

- 3 files: prv, pcf, rows
- Use paraver to analyze

❧ In the hands-on

- Add ./trace.sh to run-once.sh to enable tracing
 - > ./trace.sh \$PROGRAM ... args ...
 - > sbatch run-once.sh

Nanos6 configuration file

« Installed in share/nanos6/scripts/nanos6.toml

- [scheduler] fifo, lifo...
- [version] debug, dependencies (discrete, regions), instrument
- [cpumanager] policy (default, idle, busy, lewi, greedy)
- [throttle] enabled = false
 - tasks = 5000000
 - pressure = 70 #%
 - max_memory = "0" # half of the system memory
- [numa]
- [dlb]
- ...

Nanos++ scheduling policies

Use of schedulers:

- `fifo` `# default`
- `lifo`
- `immediate_successor` `# true by default`

Changing the policy

- `NANOS6_CONFIG_OVERRIDE="scheduler.policy = lifo"`

NANOS6 versions

« NANOS6 has runtime versions with different characteristics

- extrae
- ctf
- graph # generates the graph step by step
- lint
- stats
- verbose

« Example:

```
> NANOS6_CONFIG_OVERRIDE="instrument = extrae" \  
taskset -c 0-23 ./bin
```


Task graph generation

« Compile and link normally

```
mcc --ompss-2 -c bin.c
```

```
mcc -o bin --ompss-2 bin.o
```

« When executing specify graph as instrumentation package:

```
NANOS6_CONFIG_OVERRIDE=\
    "version.instrument=graph"
./bin
```

« Will generate one dot file per graph step

- graph-XXX-YYY.sh # script to generate PDF files
- Graph-XXX-YYY-components # directory with the graph steps

System installation

❧ Install Extrae library

- Optional, to enable tracefile generation
- Downloadable from
 - <http://tools.bsc.es/downloads>
- Binary downloads available

❧ OmpSs components downloadable from

- Nanos++ runtime
 - <https://github.com/bsc-pm/nanos6>
 - Installation: configure & make
- Mercurium compiler
 - <https://pm.bsc.es/ompss-downloads>
 - Installation: configure & make

❧ Installation guide:

- <https://pm.bsc.es/ompss-2-docs/user-guide/build/index.html>

Reporting problems

« Support mail

- pm-tools@bsc.es

« Please include snapshot of the problem

- Description
- [Mercurium] preprocessed input file (gcc -E), when possible

PATC Courses

❧ MPI + OpenMP & OmpSs

- October, 2022 [physical – at this point]

❧ Distributed programming with COMPSs

- January 25-26, 2022 [physical/online, to be decided]

❧ Performance analysis tools

- March 16-17, 2022 [physical – at this point]

❧ MPI + OmpSs, Heterogeneous prog. OpenACC/FPGAs

- March 23-25, 2022 [physical – at this point]

❧ Check <https://www.bsc.es/education/training/patc-courses>

❧ Free of charge

THANKS



**Barcelona
Supercomputing
Center**

Centro Nacional de Supercomputación

Intellectual Property Rights Notice

The User may only download, make and retain a copy of the materials for his/her use for non-commercial and research purposes.

The User may not commercially use the material, unless has been granted prior written consent by the Licensor to do so; and cannot remove, obscure or modify copyright notices, text acknowledging or other means of identification or disclaimers as they appear.

For further details, please contact BSC-CNS.

Visualizing timelines

Visualizing Paraver tracefiles

```
$ mcc --ompss -o prog -c prog.c
$ export EXTRAE_CONFIG_FILE=extrae.xml
$ NANOS6_CONFIG_OVERRIDE=\
    "version.instrumentation=extrae" taskset -c 0-7 ./prog
```

« Will generate a paraver trace

« Set of Paraver configuration files ready for OmpSs. Organized in directories:

- **Tasks: related to application tasks**
- Runtime, nanox-configs: related to OmpSs runtime internals
- **Graph_and_scheduling: related to task-graph and task scheduling**
- DataMgmt: related to data management
- CUDA: specific to GPU

Tasks' profile

« 2dp_tasks.cfg

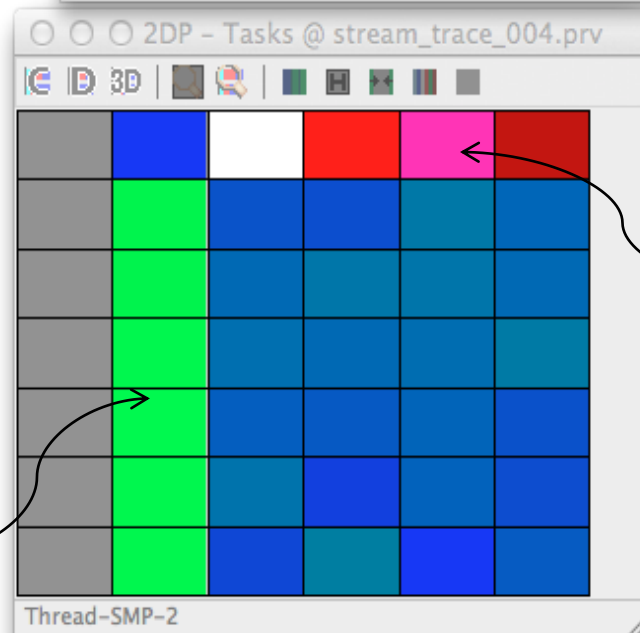

« Tasks' profile

control window:
timeline where each
color represent the
task been executed
by each thread

light blue: not executing
tasks

gradient color,
indicates given estadístic:
i.e., number of tasks instances


threads



different colours
represent different
task type

Tasks duration histogram

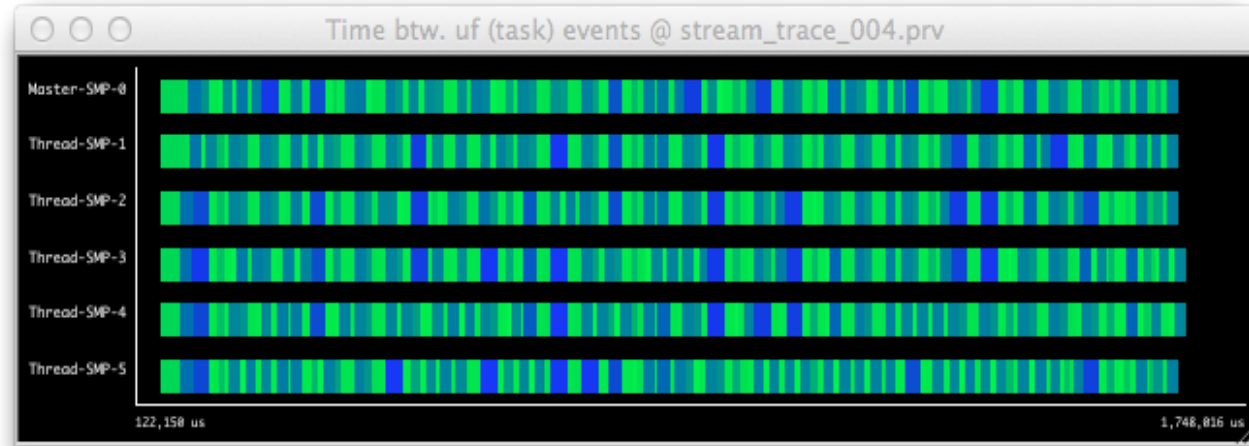
3dh_duration_task.cfg



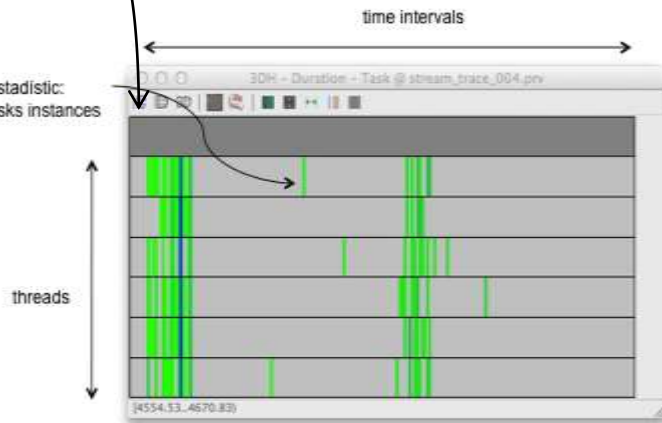
Tasks duration histogram

3dh_duration_task.cfg

control
window:
task duration



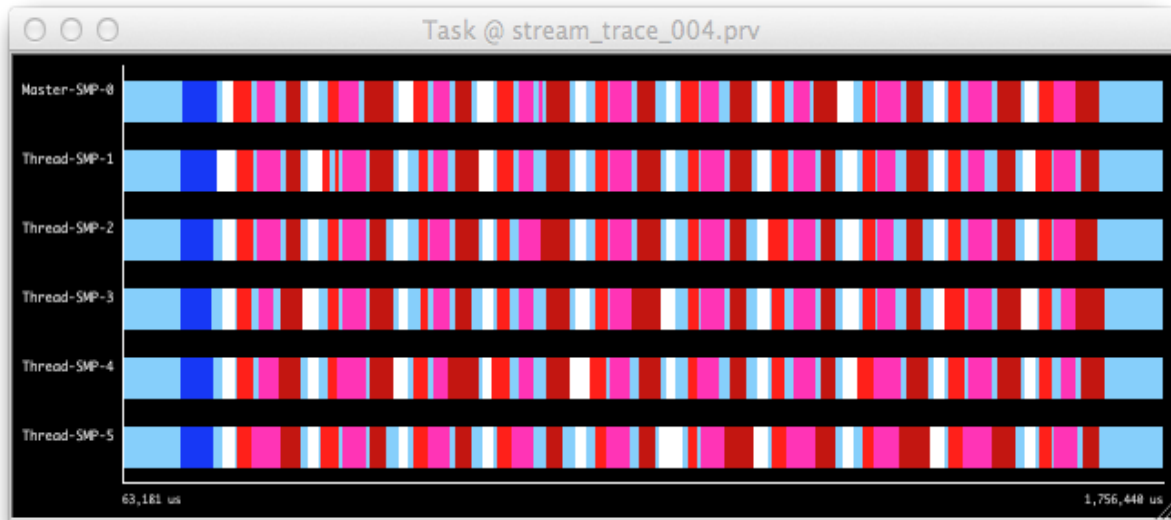
gradient color,
indicates given estadistic:
i.e., number of tasks instances



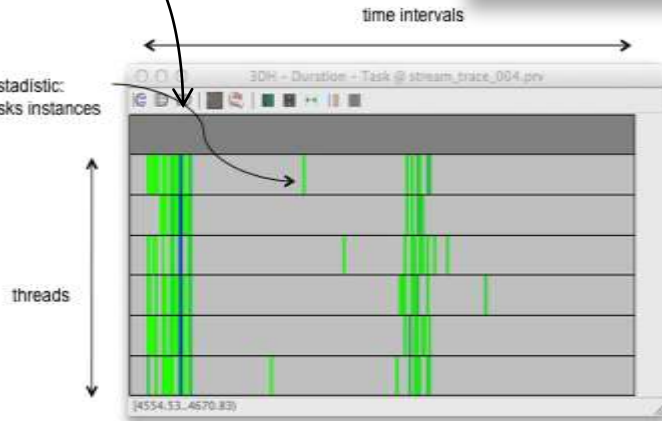
Tasks duration histogram

3dh_duration_task.cfg

3D window:
task type



gradient color,
indicates given estadistic:
i.e., number of tasks instances



Tasks duration histogram

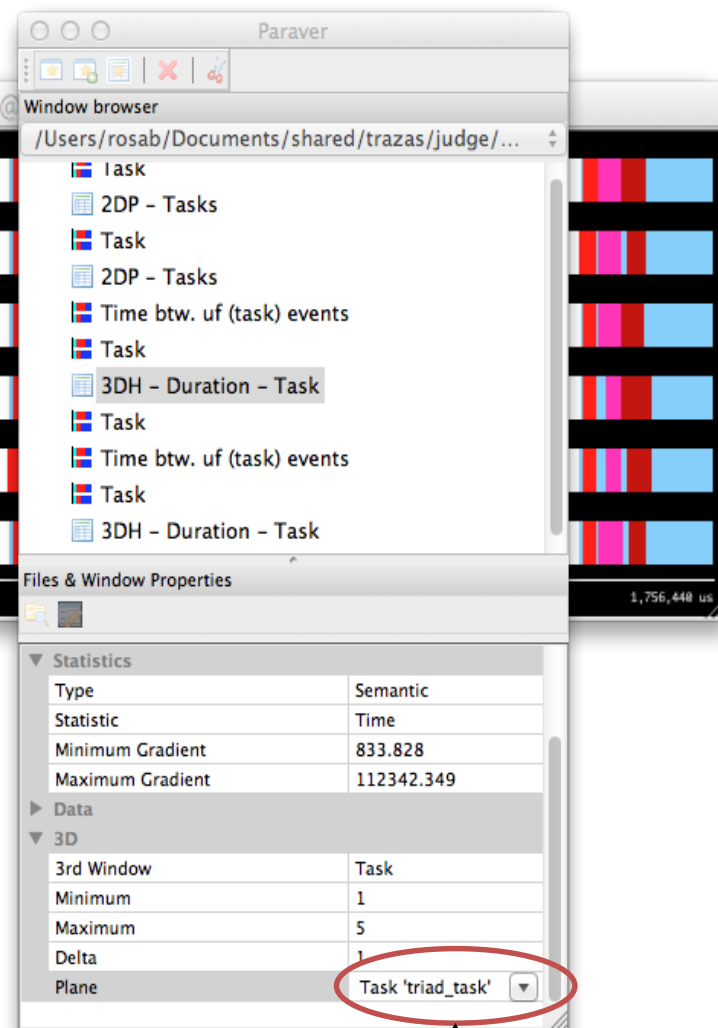

3dh_duration_task.cfg

3D window:
task type

gradient color,
indicates given estadistic:
i.e., number of tasks instances

threads

time intervals



chooser:
task type

2dp_threads_state.cfg

Thread State @ stream_trace_004.prv

Master-SMP-0

Thread-SMP-1

Thread-SMP-2

Thread-SMP-3

Thread-SMP-4

Thread-SMP-5

10,384 us

1,812,177 us

Tasks & data scoping

Inlined tasks data scope

¶ Variables for which no clause appears on the pragma:

- Arrays declared global are shared by default
- For the rest a capture value at task instantiation time semantic is followed
 - Scalars and local arrays the value is captured to initialize a private copy of the task. (== firstprivate)
 - For pointers, the pointer is captured so the task will refer to the original global object.

```
int Y[4]={1,2,3,4};

int main( )
{
    int X[4]={5,6,7,8};
    int v=0;

    #pragma oss task
    for (int i=0 ; i<4; i++) Y[i]=X[i]=v++;
    #pragma oss taskwait
        // value of v here ?
        // value of X[1] here ?
        // value of Y[1] here ?
}
```

Inlined tasks data scope

Variables for which no clause appears on the pragma:

- Arrays declared global are shared by default
- For the rest a capture value at task instantiation time semantic is followed
 - Scalars and local arrays the value is captured to initialize a private copy of the task. (== firstprivate)
 - For pointers, the pointer is captured so the task will refer to the original global object.

```
int Y[4]={1,2,3,4};

int main( )
{
    int X[4]={5,6,7,8};
    int v=0;

    #pragma oss task
    for (int i=0 ; i<4; i++) Y[i]=X[i]=v++;
    #pragma oss taskwait
        // value of v here is 0
        // value of X[1] here is 6
        // value of Y[1] here is 1
}
```


Inlined tasks data scope

Variables for which no clause appears on the pragma:

- Arrays declared global are shared by default
- For the rest a capture value at task instantiation time semantic is followed
 - Scalars and local arrays the value is captured to initialize a private copy of the task. (== firstprivate)
 - For pointers, the pointer is captured so the task will refer to the original global object.

```
int main( )
{
    int *X;
    int v=0;

    X = (int *) malloc(4*sizeof(int));

    #pragma oss task
    for (int i=0 ; i< 4; i++) X[i]=v++;
    #pragma oss taskwait

    //value of X[1] here ?
}
```

Inlined tasks data scope

Variables for which no clause appears on the pragma:

- Arrays declared global are shared by default
- For the rest a capture value at task instantiation time semantic is followed
 - Scalars and local arrays the value is captured to initialize a private copy of the task. (== firstprivate)
 - For pointers, the pointer is captured so the task will refer to the original global object.

```
int main( )
{
    int *X;
    int v=0;

    X = (int *) malloc(4*sizeof(int));

    #pragma oss task
    for (int i=0 ; i< 4; i++) X[i]=v++;
    #pragma oss taskwait

    //value of X[1] here is 1
}
```

Inlined tasks data scope

❧ Data scope (... as in OpenMP)

- Variables can be explicitly declared as :
 - Shared
 - Private
 - Firstprivate

❧ In case of doubt

- Default (none)
- Explicitly declare

```
int Y[4]={1,2,3,4};
int main()
{
    int X[4]={5,6,7,8};
    int v=0;

    #pragma omp task shared (v, X)
    for (int i=0 ; i<4; i++) Y[i]=X[i]=v++;
    #pragma omp taskwait

    // value of v here ?
    // value of X[1] here ?
    // value of Y[1] here ?
}
```

Inlined tasks data scope

❧ Data scope (... as in OpenMP)

- Variables can be explicitly declared as :
 - Shared
 - Private
 - Firstprivate

❧ In case of doubt

- Default (none)
- Explicitly declare

```
int Y[4]={1,2,3,4};
int main()
{
    int X[4]={5,6,7,8};
    int v=0;

    #pragma oss task shared (v, X)
    for (int i=0 ; i<4; i++) Y[i]=X[i]=v++;
    #pragma oss taskwait

    // value of v here is 4
    // value of X[1] here is 1
    // value of Y[1] here is 1
}
```

Inlined tasks data scope

Explicit scope declaration

- Firstprivate: a private instance of the data initialized to the value of the original variable at task instantiation time.
- Private: an uninitialized instance is allocated and used by the task
- In both cases, the private data instance is deleted on task completion.

```
int X[4]={1,2,3,4};

int main( )
{
int i=2;

#pragma oss task firstprivate (i)
for ( ; i< 4; i++) X[i]=i;
#pragma oss taskwait

    // value of i here ?
    // value of X[0] here ?
    // value of X[3] here ?
}
```

```
int main( )
{
int X[4]={1,2,3,4};

int i=2;

#pragma oss task private(i) shared(X)
for (i=0; i< 100; i++) X[i]=i;
#pragma oss taskwait

    // value of i here ?
    // value of X[0] here ?
    // value of X[3] here ?
}
```

Inlined tasks data scope

Explicit scope declaration

- Firstprivate: a private instance of the data initialized to the value of the original variable at task instantiation time.
- Private: an uninitialized instance is allocated and used by the task
- In both cases, the private data instance is deleted on task completion.

```
int X[4]={1,2,3,4};

int main( )
{
int i=2;

#pragma oss task firstprivate (i)
for ( ; i< 4; i++) X[i]=i;
#pragma oss taskwait

    // value of i here is 2
    // value of X[0] here is 1
    // value of X[3] here is 3
}
```

```
int main( )
{
int X[4]={1,2,3,4};

int i=2;

#pragma oss task private(i) shared(X)
for (i=0; i< 100; i++) X[i]=i;
#pragma oss taskwait

    // value of i here is 2
    // value of X[0] here is 0
    // value of X[3] here is 3
}
```


Additional directives

Tasks: the if clause

`#pragma omp task [if(...)]`

- if (expr): If expression evaluates to false, task will be created but will be executed immediately (not deferred)
 - Scheduling rather than overhead
 - OpenMP issues if taskwait inside the structured block and it is executed inline as part of the calling region

```
void foo(int *size, int **vector, int N) {  
    for (int i=0; i<N; i++) {  
        #pragma omp task if ( size[i] > MIN_SIZE )  
        compute(vector[i], size[i]) ;  
    }  
}
```



Outlined directives

Outlined tasks instantiation

Task pragma attached to function definition

- All function invocations become a task
- The programmer gives a name, this enables later to provide several implementations

```
#pragma oss task
void foo (int Y[size], int size) {
    for (int j=0; j<size; j++) Y[j]= j;
}

int Y[4]={1,2,3,4};

int main()
{
    int X[100];

    foo (X, 100) ;
    #pragma oss taskwait
    ...
}
```


Invocation
instantiates task

foo

Outlined tasks data scoping

¶ The semantic is capture value at task instantiation time (function invocation)

- For scalars, the value is captured. Equivalent to `firstprivate`
- For pointers, the value of the pointer is captured
- For arrays, the address is captured



```
#pragma oss task
void foo (int Y[size], int size, int a)
{
    for (int j=0; j<size; j++) Y[j]= a++;
}
```

```
int Y[6]={1,2,3,4,5,6};

int main()
{
    int v=0;
    int X[4]={5,6,7,8};

    int *Z;
    Z=(int *)malloc(4*sizeof(int));


    foo (X, 4, v) ;
    foo (Y, 2, v+2);
    foo (&Y[3], 3, v++);
    foo (Z,4,v);
    #pragma oss taskwait

    // value of v here ?
    // value of X[1] here ?
    // value of Y[4] here ?
    // value of Y[3] here ?
    // value of Z[1] here ?
}
```

Outlined tasks data scoping

¶ The semantic is capture value at task instantiation time (function invocation)

- For scalars, the value is captured. Equivalent to `firstprivate`
- For pointers, the value of the pointer is captured
- For arrays, the address is captured



```
#pragma oss task
void foo (int Y[size], int size, int a)
{
    for (int j=0; j<size; j++) Y[j]= a++;
}
```

```
int Y[6]={1,2,3,4,5,6};

int main()
{
    int v=0;
    int X[4]={5,6,7,8};

    int *Z;
    Z=(int *)malloc(4*sizeof(int));

    foo (X, 4, v) ;
    foo (Y, 2, v+2);
    foo (&Y[3], 3, v++);
    foo (Z,4,v);
    #pragma oss taskwait

    // value of v here is 1
    // value of X[1] here is 1
    // value of Y[4] here is 1
    // value of Y[3] here is 0
    // value of Z[1] here is 2
}
```

Defining dependences for outlined tasks

« Clauses that express data direction:


- Input, output, inout
- The argument is an lvalue expression based on data visible at the point of declaration (global variables and arguments)
- The object pointed by the lvalue expression will be used to compute dependences.

```
#pragma oss task out(*px)
void set (int *px, int v) {*px = v;}

#pragma oss task inout(*px)
void incr (int *px) {(*px)++;}

#pragma oss task in(x)
void do_print (int x) {
    printf("from do_print %d\n" , x );
}
```

```
set(&x,5);      //1
do_print(x);    //2
incr(&x);       //3
do_print(x);    //4
#pragma oss taskwait
```



Defining dependences for outlined tasks

« Clauses that express data direction:


- Input, output, inout
- The argument is an lvalue expression based on data visible at the point of declaration (global variables and arguments)
- The object pointed by the lvalue expression will be used to compute dependences.

```
#pragma oss task out(*px)
void set (int *px, int v) {*px = v;}

#pragma oss task in(x)
int incr (int x) {int y; y= x+1; return(y);}

#pragma oss task in(x)
void do_print (int x) {
    printf("from do_print %d\n" , x );
}
```

```
set(&x,5);          //1
do_print(x);        //2
x = incr(x);         //3
do_print(x);        //4
#pragma oss taskwait
```



Mixing inlined and outlined tasks


```
#pragma oss task in (x)
void do_print (int x) {
    printf("from do_print %d\n" , x ) ;
}

int main()
{
    int x;

    x=3;

    #pragma oss task out( x )
    x = 5; //1
    #pragma oss task in( x )
    printf("from main %d\n" , x ); //2
    do_print(x); //3
    #pragma oss task inout( x )
    x++; //4
    #pragma oss task in( x )
    printf ("from main %d\n" , x ); //5
}
```

non-taskified:
executed
sequentially



Partial control flow synchronization

`#pragma oss taskwait on (lvalue_expr-list)`

- Expressions allowed are the same as for the dependency clauses
- Blocks the encountering task until the data is available


```
#pragma oss task in([N][N]A, [N][N]B) inout([N][N]C)
void dgemm(float *A, float *B, float *C);

main() {
(
  dgemm(A,B,C); //1
  dgemm(D,E,F); //2
  dgemm(C,F,G); //3
  dgemm(A,D,H); //4
  dgemm(C,H,I); //5


  #pragma oss taskwait on (F)
  printf("result F = %f\n", F[0][0]);

  dgemm(H,C,J); //6

  #pragma oss taskwait
  printf("result J = %f\n", J[0][0]);
}
```



Array sections in outlined tasks



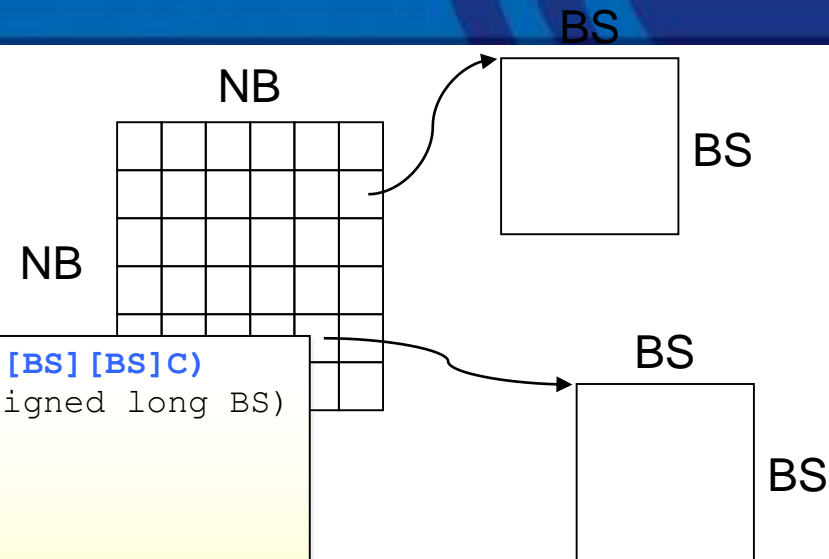
```
#pragma oss task in([n]vec) inout([n]c)
void sum_task ( int *vec , int n , int *c);

void main(){
    int actual_size;

    for (int j; j<N; j+=BS){
        actual_size = (N- j > BS ? BS: N-j);
        sum_task (&vec[j], actual_size, &c[j]);
    }
}
```

dynamic size of
argument

Array sections in outlined tasks



```
#pragma oss task in([BS][BS]A, [BS][BS]B) inout([BS][BS]C)
void matmul(double *A, double *B, double *C, unsigned long BS)
{
    int i, j, k;

    for (i=0; i<BS; i++)
        for (j=0; j<BS; j++)
            for (k=0; k<BS; k++)
                C[i*BS+j] += A[i*BS+k] * B[k*BS+j];
}
```

```
void compute(unsigned long NB, unsigned long BS,
             double *A[NB][NB], double *B[NB][NB], double *C[NB][NB])
{
    unsigned i, j, k;


    for (i=0; i<NB; i++)
        for (j=0; j<NB; j++)
            for (k=0; k<NB; k++)
                matmul (A[i][k], B[k][j], C[i][j], BS);
}
```

Array sections in outlined tasks

N = total size of matrix

BS = block size

NB = number of blocks



```
#pragma oss task in([BS][BS]A, [BS][BS]B) inout([BS][BS]C)
void matmul(double *A, double *B, double *C, unsigned long BS)
{
    int i, j, k;

    for (i=0; i<BS; i++)
        for (j=0; j<BS; j++)
            for (k=0; k<BS; k++)
                C[i*BS+j] += A[i*BS+k] * B[k*BS+j];
}
```

```
void compute(unsigned long BS, unsigned long N, double *A, double *B, double *C)
{
    unsigned i, j, k;

    for (i = 0; i < N; i+=BS)
        for (j = 0; j < N; j+=BS)
            for (k = 0; k < N; k+=BS)
                matmul (&A[i*N+k*BS], &B[k*N+j*BS], &C[i*N+j*BS], BS);
}
```


Specification of incomplete data-directionality

```
#pragma oss task out (*sentinel)
void foo ( .... , int *sentinel){
    ...
}

#pragma oss task in (*sentinel)
void bar ( .... , int *sentinel){
    ...
}

main () {
    int sentinel;

    foo (... , &sentinel);
    bar (... , &sentinel)
}
```



- Mechanism to handle complex dependences
 - when difficult to specify proper input/output clauses
- To be avoided if possible
 - the use of an element or group of elements as sentinels to represent a larger data-structure is valid
 - however might made code non-portable to heterogeneous platforms if `copy_in/out` clauses cannot properly specify the address space that should be accessible in the devices

Specification of incomplete data-directionality

- ❧ Directionality not required for all arguments
- ❧ May even be used with variables not accessed in that way or even used
 - used to force dependences under complex structures (graphs, ...)

```
#pragma oss task in(*A, *B) inout(*C)
void matmul(double *A, double *B, double *C,
            unsigned long BS)
{
    int i, j, k;

    for (i = 0; i < BS; i++)
        for (j = 0; j < BS; j++)
            for (k = 0; k < BS; k++)
                C[i][j] += A[i*BS+k]*B[k*BS+j];
}
```

Using element C[0][0] as representative/sentinel for the whole block.

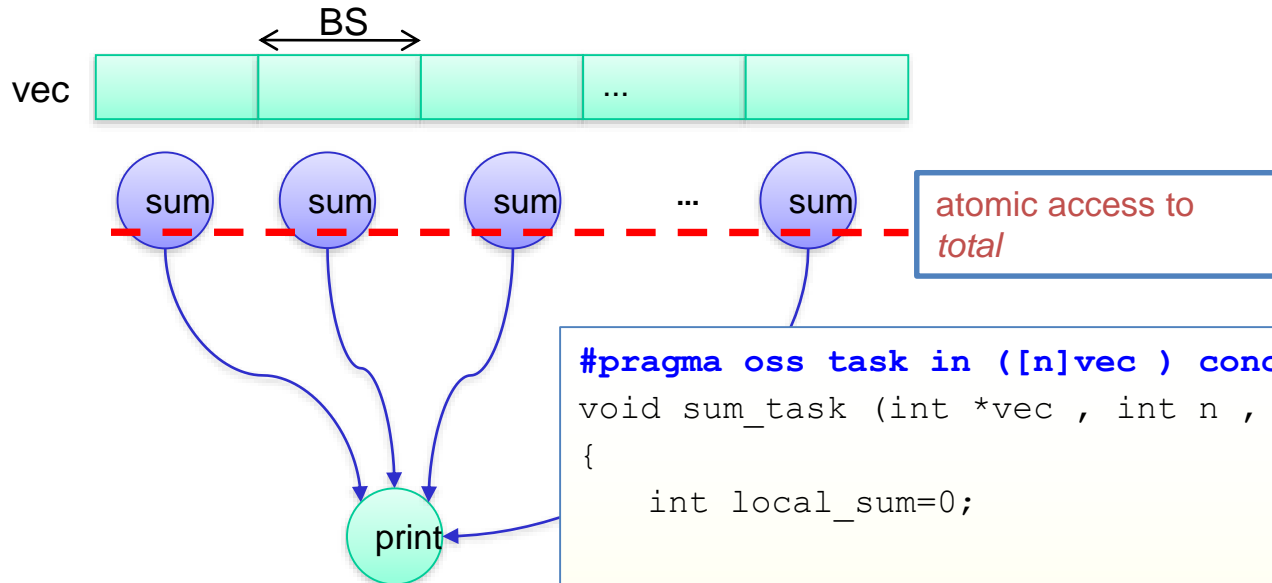
Will build proper dependences between tasks.

Does NOT provide actual information of data access pattern. (see copy clauses)

- Pragmas outlined

```
program example
parameter(N=2048)
integer, parameter :: BSIZE = 64
real v(N), vx(N), vy(N), vz(N)
integer :: jj
...
interface
    !$oss task out (v(1:BSIZE)) in (vx(1:BSIZE), vy(1:BSIZE), vz(1:BSIZE)) \
    label(vmod)
    subroutine v_mod(BSIZE, v, vx, vy, vz)
        implicit none
        integer, intent(in) :: BSIZE
        real, intent(out) :: v(BSIZE)
        real, intent(in), dimension(BSIZE) :: vx, vy, vz
    end subroutine
end interface
...
do jj=1, N, BSIZE
call v_mod(BSIZE, v(jj), vx(jj), vy(jj), vz(jj))
enddo
...
!$oss taskwait
```

Concurrent in outlined tasks



```
#pragma oss task in ([n]vec ) concurrent (*result)
void sum_task (int *vec , int n , int *result)
{
    int local_sum=0;

    for (int i = 0; i < n ; i ++ )
        local_sum += vec [i] ;

    #pragma oss atomic
    *result += local_sum;
}

void main() {
    for (int j=0; j<N; j+=BS)
        sum_task (&vec[j], BS, &result);
    #pragma oss task in (result)
    printf ("TOTAL is %d\n", result);
}
```

Commutative in outlined tasks

